

7.6 Problem set 1: The prisoners' dilemma

7.6.1 Background

I'm an apple farmer who hates apples but loves broccoli. You're a broccoli farmer who hates broccoli but loves apples. The obvious solution to this sad state of affairs is for us to trade – I ship you a box of my apples and you ship me a box of your broccoli. Win-win.

But I might try to get clever by shipping an empty box. Instead of cooperating, I “defect”. I still get my broccoli from you (assuming you don't defect) and get to keep my apples. You, thinking through this scenario, realize that you're better off defecting as well; at least you'll get to keep your broccoli. But then, nobody gets what we want; we're both worse off. The best thing to do in this **DONATION GAME** seems to be to defect.

It's a bit of a mystery, then, why people cooperate at all. The answer may lie in the fact that we engage in many rounds of the game. If you get a reputation for cooperating, others may be willing to cooperate as well, leading to overall better outcomes for all involved.

The donation game is an instance of a classic game-theory thought experiment called the **PRISONER'S DILEMMA**. A prisoner's dilemma is a type of game involving two players in which each player is individually incentivized to choose a particular action, even though it may not result in the best global outcome for both players. The outcomes are commonly specified through a payoff matrix, such as the one in Table 7.2.

To read the matrix, Player 1's actions are outlined at the left and Player 2's actions at the top. The entry in each box corresponds to a payoff to each player, depending on their respective actions. For instance, the top-right box indicates the payoff when Player 1 cooperates and Player 2 defects. Player 1 receives a payoff of -2 and Player 2 receives a payoff of 5 in that case.

To see why a dilemma arises, consider the possible actions taken by Player 1. If Player 2 cooperates, then Player 1 should defect rather than cooperating, since the payoff from defecting is higher ($5 > 3$). If Player 2 defects, then Player 1 should again defect since the payoff from defecting is higher ($5 > -2$). The same analysis applies to Player

Table 7.1: Comparison of three structuring mechanisms: tuples, records, and lists.

		Player 2	
		Cooperate	Defect
Player 1	Cooperate	(3, 3)	(-2 , 5)
	Defect	(5, -2)	(0, 0)

Table 7.2: Example payoff matrix for a prisoner's dilemma. This particular payoff matrix corresponds to a donation game in which providing the donation (of apples or broccoli, say) costs 2 unit and receiving the donation provides a benefit of 5 units.

2. Therefore, both players are incentivized to defect. However, the payoff from both players defecting (each getting 0) is objectively worse for both players than the payoff from both players cooperating (each getting 4).

An **ITERATED PRISONER'S DILEMMA** is a multi-round prisoner's dilemma, where the number of rounds is not known.¹⁴ A **STRATEGY** specifies what action to take based on a history of past rounds of a game. We can (and will) represent a history as a list of pairs of actions (cooperate or defect) taken in the past, and a strategy as a function from histories to actions.

For example, a simple strategy is to ignore the histories and always defect. We call that the “nasty” strategy. More optimistically is the “patsy” strategy, which always cooperates.

Whereas the above analysis showed both players are incentivized to defect in a single-round prisoner's dilemma (leading to the nasty strategy), that is no longer necessarily the case if there are multiple rounds. Instead, more complicated strategies can emerge as players can take into account the history of their opponent's plays and their own. A particularly effective strategy – effective because it leads to cooperation, with its larger payoffs – is **TIT-FOR-TAT**. In the tit-for-tat strategy, the player starts off by cooperating in the first round, and then in later rounds chooses the action that the other player just played, rewarding the other player's cooperation by cooperating and punishing the other player's defection by defecting.

In this problem set, you'll complete a simulation of the iterated prisoner's dilemma that allows for testing different payoff matrices and strategies.

7.6.2 Setup

We'll assume that you have successfully configured your development environment as specified by Problem Set 0. If not, you'll want to do that before getting started.

To download the assignment, accept the assignment [here](#) and follow any instructions. Open the directory 1, and you will find a file (`ps1.ml`) that contains detailed instructions and scaffolding for all of the problems. In particular, it provides a stub for each function you are asked to implement. Typically, that stub merely raises an exception noting that the function has not yet been written – something like¹⁵

```
let merge =
  fun _ -> failwith "merge not implemented" ;;
```

You'll want to replace these stubs with a correct implementation of

¹⁴ If the number of rounds is known by the players ahead of time, players are again incentivized to defect for all rounds. We will not delve into the reasoning here, as that is outside the scope of this course, but it is an interesting result!

¹⁵ The `failwith` function will be explained in Section 10.3. For the time being, you can treat this as a fixed idiom.

the function. Other parts of the assignment may have you perform other tasks to check your understanding of the material.

Here are some things to keep in mind as you complete the assignments.

Style Good style is an important part of programming well. It is useful not only in making code more readable but also in helping identify syntactic and type errors in your functions. For this and all assignments, your code should adhere to the style guide in Chapter C.

Compilation errors In order to submit your work to the course grading server, *your solution must compile against our test suite*. The system will reject submissions that do not compile. If there are problems that you are unable to solve, you must still write a function that matches the expected type signature, or your code will not compile. When we provide stub code, that code will typically compile to begin with. If you are having difficulty getting your code to compile, please visit office hours or post on Piazza. *Emailing your homework to your TF or the Head TFs is not a valid substitute for submitting to the course grading server. Please start early.*

Helper functions Feel free to use helper functions to make your code cleaner, more modular, and easier to test.

7.6.3 Some practice functions

To get started, you'll write a series of functions that perform simple manipulations over lists, strings, numbers, and booleans. (Some of these may be useful later in implementing the iterated prisoner's dilemma.) See the comments in `ps1.ml` for the specifications. Give the functions the names listed in the comments, as they must be named as specified in order to compile against our automated unit tests.

The best way to learn about the core language is to work directly with the core language features. Consequently, you should not use any library functions in implementing your solutions to this problem set.

7.6.4 Unit testing

Thorough testing is important in all your work, and we hope to impart this view to you in CS51. Testing will help you find bugs, avoid mistakes, and teach you the value of short, clear functions. In the file `ps1_tests.ml`, we've put some prewritten tests for Problem 1a using

the testing method of Section 6.5. Spend some time understanding how the testing function works and why these tests are comprehensive. Then, for each function in Problem 1, write tests that thoroughly test the functionality of each of the remaining sections of Problem 1, covering all code paths and corner cases.

To run your tests, run the shell command

```
% make ps1_tests.byte
```

in the directory where your `ps1.ml` is located and then run the compiled program by executing `ps1_tests.byte`:

```
% ./ps1_tests.byte
```

The program should then generate a full report with all of the unit tests for all of the functions in the problem set. *Remember: your assignment is not complete if you have not written unit tests for your work.*

7.6.5 The prisoner's dilemma

In this section, you will apply functional programming concepts to complete a model of the **iterated prisoner's dilemma**, including the implementation of various strategies, including one of your own devising.

Follow the comments in `ps1.ml` for the specifications. Feel free to use any of the functions that you implemented in the earlier parts of the problem set.

All of the programming concepts needed to do the problem set have already been introduced. There's no need to use later constructs, and you should refrain from doing so. In particular, you'll want to avoid imperative programming, and you should not use any library functions.

7.6.6 Tournament

To allow you to see how your custom strategy fares, we will run a class-wide round-robin tournament in which your strategy will be run against every other student's strategy for a random number of rounds. To keep things interesting, we will add a small amount of noise to each strategy: 5% of the time, we will use the opposite action of the one a strategy specifies.

The tournament is pseudonymized and optional. If you want to participate, you'll provide a pseudonym for your entrant. If you specify an empty pseudonym, we won't enter you in the tournament, though we encourage you to come up with clever strategies to compete against your fellow classmates! We'll post and regularly update

a leaderboard, displaying each student's pseudonym and their strategy's average payoff. The top score will win a rare and wonderful prize!

7.6.7 Testing

For this section, you should again provide unit tests in `ps1_tests.ml` for each function that you implement in Problem 2.

In addition, you can choose to uncomment the `main` function in `ps1.ml` and then recompile the file by running `make ps1.byte` in your shell. Then, run the command `./ps1.byte` and you should see via printed output that Nasty earns a payoff of 500 and Patsy earns a payoff of -200. You can change the strategies played to see how different strategies perform against each other.

7.6.8 Troubleshooting

We provide a file `ps1.mli` that checks to make sure your functions in `ps1.ml` have the correct types. (We will discuss `.mli` files and the related concept of module signatures in some detail in Chapter 12.) You *should not* modify `ps1.mli`.

If you change the type of a function, you may see an error like:

```
Error: The implementation ps1.ml does not match the
      interface ps1.cmi:
      Values do not match:
        val nonincreasing : int list -> int
      is not included in
        val nonincreasing : int list -> bool
```

This error means that your function `nonincreasing` has the wrong type. You should look back at your code and modify it so that it has the type specified in the assignment.

If you delete a function, you may see an error like:

```
Error: The implementation ps1.ml does not match the
      interface ps1.cmi:
      The value `nonincreasing` is required but not provided
```

This error means that you are missing an implementation of the function `nonincreasing`. If you are not able to get a particular function to compile, do not delete it. Instead, you can revert the function back to the stub implementation in the distribution code. (You can look at previous commits on GitHub to see what the stub implementation was.)

7.6.9 Submission

Before submitting, please estimate how much time you spent on the assignment by editing the line:

```
let minutes_spent_on_pset () : int = failwith "not provided" ;;
```

to replace the value of the function with an approximate estimate of how long (in minutes) the assignment took you to complete. For example, if you spent 6 hours on this assignment, you should change the line to:

```
let minutes_spent_on_pset () : int = 360 ;;
```

We also ask that you reflect on the process you went through in working through the problem set, where you ran into problems and how you ended up resolving them. What might you have done in retrospect that would have allowed you to generate as good a submission in less time? Please provide your thoughts by editing the value of the `reflect` function. It should look something like

```
let reflection () : string =  
  "...your reflections go here..." ;;
```

Make sure your code still compiles. Then, to submit the assignment, follow the instructions found in Section [A.4.3](#).